

## 1. Hyper-parameter optimization

The hyper-parameters for the MLP, CNN and RN models were determined through random search, in which a model is repeatedly trained with randomly sampled parameters and evaluated on the validation set. This process was repeated 400 times for all MLP and CNN models, and 200 times for the RN models. The space of possible parameters that were sampled from is listed in Table 1 for the MLP models, Table 2 for the CNN models and Table 3 for the RN models. The architecture and parameters of the models that had the best AUC on the validation set are illustrated in Section 2.

<b>Feature extractor</b>	
Number of dense layers	[1, 2]
Neurons	[25, 50, 100, 200]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Batch normalization	[true, false]
Activity regularization	[1e-2, 1e-3, 1e-4, 1e-5, 0]
Kernel regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
Bias regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
<b>Combiner</b>	
Number of dense layers	1
Neurons	[6, 10, 20]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Batch normalization	[true, false]
Activity regularization	[1e-2, 1e-3, 1e-4, 1e-5, 0]
Kernel regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
Bias regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
<b>Classifier</b>	
Number of dense layers	1
Neurons	10
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Batch normalization	[true, false]
Activity regularization	[1e-2, 1e-3, 1e-4, 1e-5, 0]
Kernel regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
Bias regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
<b>Additional settings</b>	
Optimizer	Adam
Learning-rate	[3e-3, 1e-3, 3e-4, 1e-4, 3e-5, 1e-5]
Epochs	100
Batch-size	64

Table 1: MLP hyper-parameter search space. The dense layer in the “Classifier” part was only included for models using additional clinical variables. All parameters were sampled uniformly at random from the options shown in the brackets.

<b>Feature extractor</b>	
Number of conv layers	[2, 3, 4]
Pool-size	[[7-30], [4-10], [3-6]]
Filters, first	[8-64]
Filters, last	[8-64]
Kernel size, first	[5-65]
Kernel size, last	[5-65]
Batch normalization	[true, false]
Weight decay	[0.1, 0.01, 0.001, 0]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Number of dense layers	[0, 1]
Neurons	[10, 50, 100]
Dense dropout	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
<b>Combiner</b>	
Number of dense layers	[0, 1]
Neurons	[10, 20]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
<b>Classifier</b>	
Number of dense layers	1
Neurons	[10, 20, 50, 100]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
<b>Additional settings</b>	
Optimizer	Adam
Learning-rate	[3e-3, 1e-3, 3e-4, 1e-4, 3e-5, 1e-5]
Epochs	100
Batch-size	64

Table 2: CNN hyper-parameter search space. The “Feature extractor” part begins with 2, 3 or 4 convolutional blocks, and the pool-size in each case was sampled from the range [7-30], [4-10] or [3-6] respectively. The kernel size and number of filters for each convolutional block was linearly interpolated between the first and last values indicated in the table. Each convolutional block included a potential batch-normalization, followed by max-pooling and dropout. After all the convolutions there was a flatten layer, potentially followed by a dense-layer. The dense layer in the “Combiner” part was only used for the models that included additional clinical variables. All parameters were sampled uniformly at random from the options shown in the brackets.

<b>Feature extractor</b>	
Number of dense layers	[1, 2]
Neurons	[25, 50, 100, 200]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Batch normalization	[true, false]
Activity regularization	[1e-2, 1e-3, 1e-4, 1e-5, 0]
Kernel regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
Bias regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
<b>Combiner</b>	
Number of dense layers	1
Neurons	[6, 10, 20]
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Batch normalization	[true, false]
Activity regularization	[1e-2, 1e-3, 1e-4, 1e-5, 0]
Kernel regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
Bias regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
<b>Classifier</b>	
Number of dense layers	1
Neurons	10
Dropouts	[0.0, 0.1, 0.2, 0.3, 0.4, 0.5]
Batch normalization	[true, false]
Activity regularization	[1e-2, 1e-3, 1e-4, 1e-5, 0]
Kernel regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
Bias regularization	[1e-1, 1e-2, 1e-3, 1e-4, 0]
<b>Additional settings</b>	
Optimizer	Adam
Learning-rate	[3e-3, 1e-3, 3e-4, 1e-4, 3e-5, 1e-5]
Epochs	50
Batch-size	32

Table 3: RN hyper-parameter search space. The dense layer in the “Classifier” part was only included for models using additional clinical variables. Two learning-rates were sampled for each model, and the higher rate was used for the first epochs (a random number between 5 and 25), after which the lower learning rate was used. All parameters were sampled uniformly at random.

## 2. Final model parameters

### 2.1. Logistic regression models




Figure 1: Parameters and structure of the logistic regression model using only the index ECG as input.

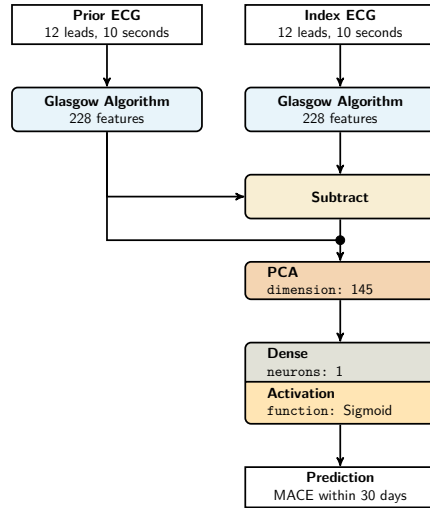


Figure 2: Parameters and structure of the logistic regression model using both the index and prior ECGs as input. Subtraction is vector-valued. The circle denotes concatenation.

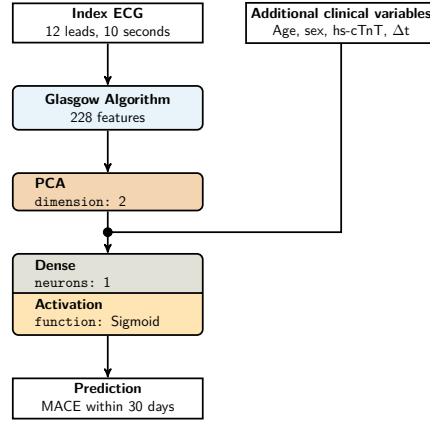


Figure 3: Parameters and structure of the logistic regression model using index ECG and additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation.

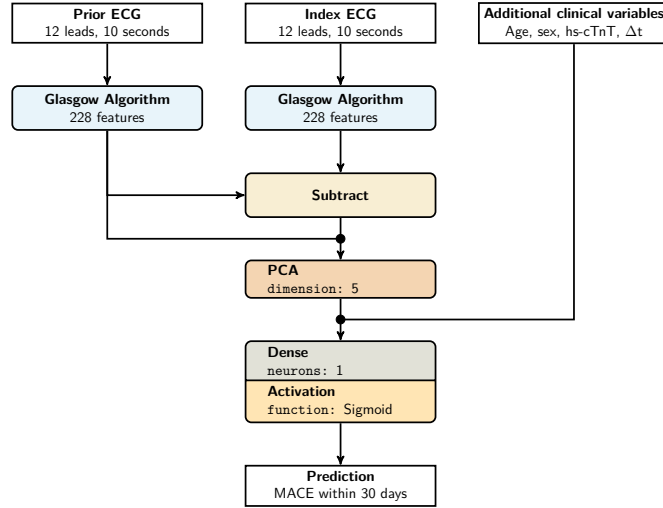


Figure 4: Parameters and structure of the logistic regression model using both index and prior ECGs together with additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. Subtraction is vector-valued. The circle denotes concatenation.

## 2.2. Multilayer perceptron models




Figure 5: Parameters and structure of the multilayer perceptron model using only the index ECG as input. The learning rate was 0.0001.

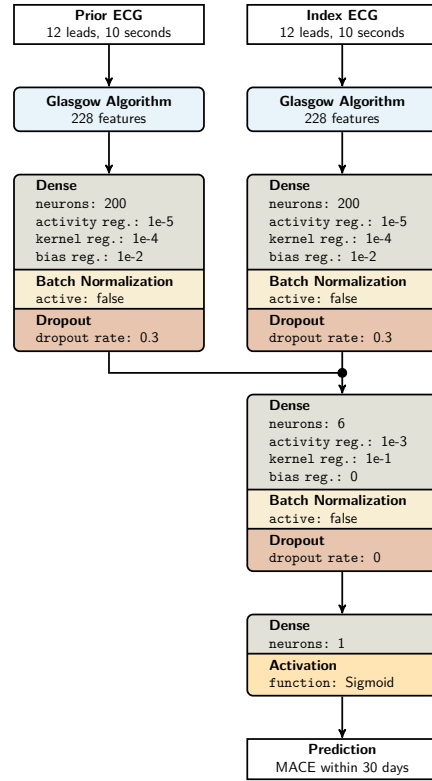


Figure 6: Parameters and structure of the multilayer perceptron model using both the index and prior ECGs as input. The circle denotes concatenation. The learning rate was 0.0001.

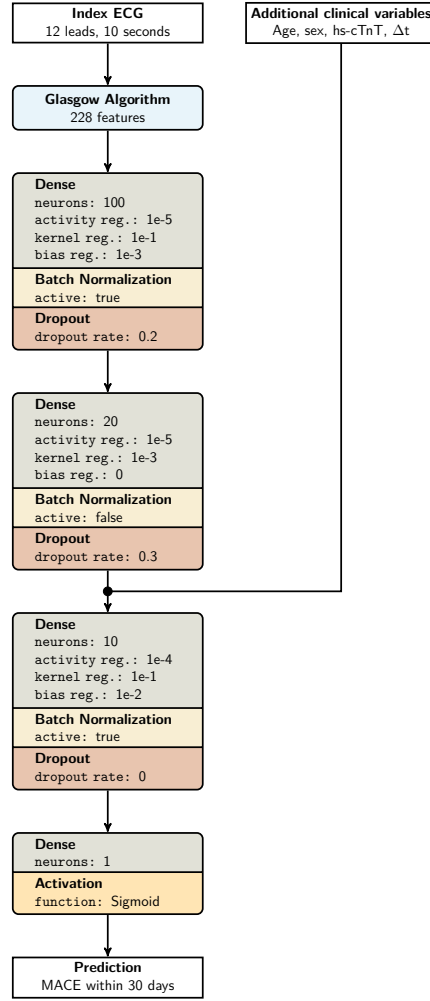


Figure 7: Parameters and structure of the multilayer perceptron model using index ECG and additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation. The learning rate was 0.0003.






Figure 8: Parameters and structure of the multilayer perceptron model using both index and prior ECGs together with additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation. The learning rate was 0.001.

### 2.3. Convolutional neural network models




Figure 9: Parameters and structure of the convolutional neural network model using only the index ECG as input. The learning rate was 0.0001.

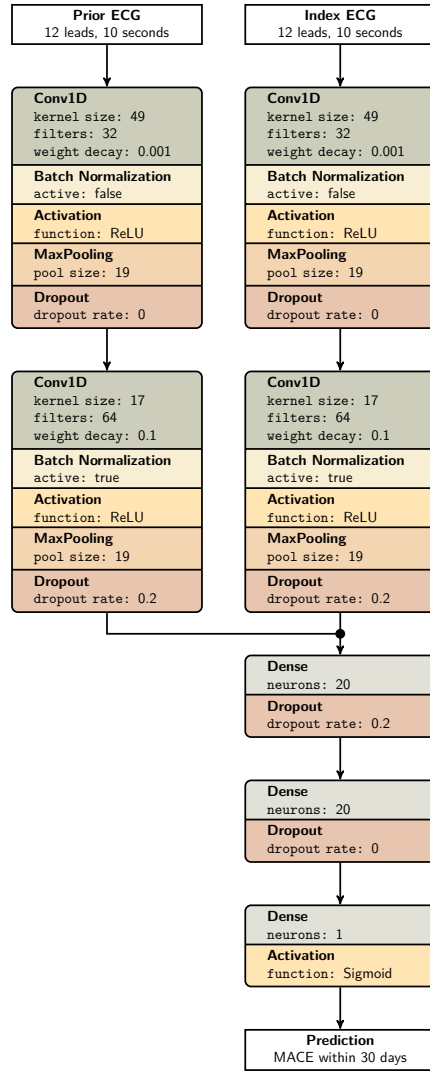


Figure 10: Parameters and structure of the convolutional neural network model using both the index and prior ECGs as input. The circle denotes concatenation. The learning rate was 0.00001.

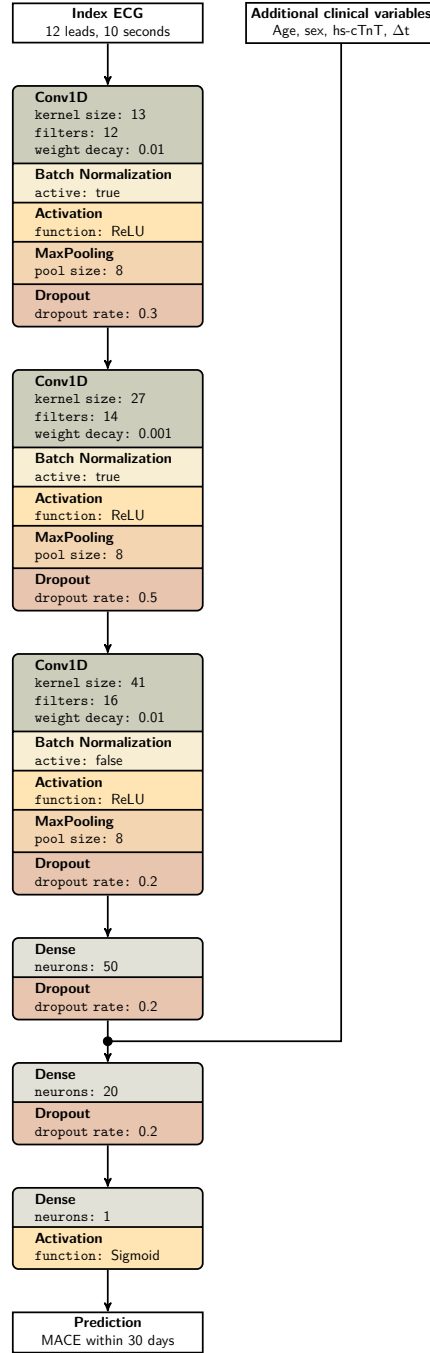


Figure 11: Parameters and structure of the convolutional neural network model using index ECG and additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation. The learning rate was 0.001.




Figure 12: Parameters and structure of the convolutional neural network model using both index and prior ECGs together with additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation. The learning rate was 0.001.

## 2.4. Residual neural network models

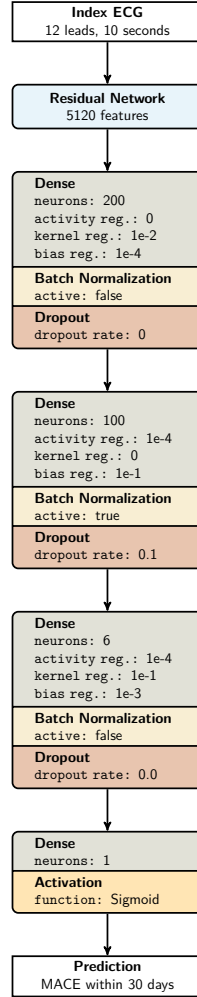


Figure 13: Parameters and structure of the residual network model using only the index ECG as input. The learning rate was 0.0003 for all 50 epochs.

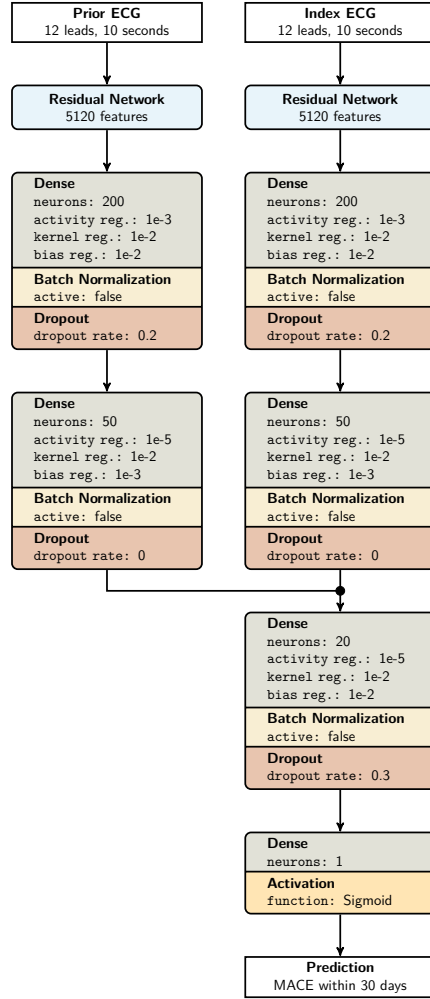


Figure 14: Parameters and structure of the residual network model using both the index and prior ECGs as input. The circle denotes concatenation. The learning rate was 0.001 for the first 18 epochs, and 0.00001 for the final 32 epochs.

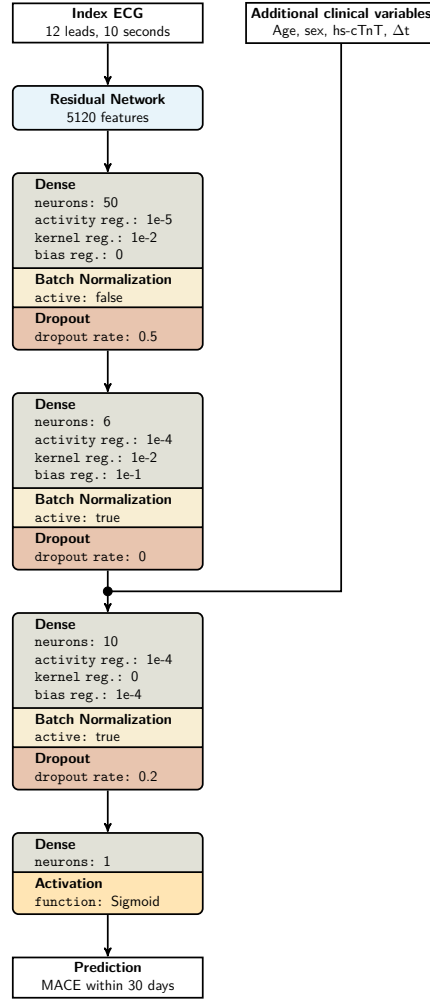


Figure 15: Parameters and structure of the residual network model using index ECG and additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation. The learning rate was 0.0003 for the first 24 epochs, and 0.00003 for the final 26 epochs.



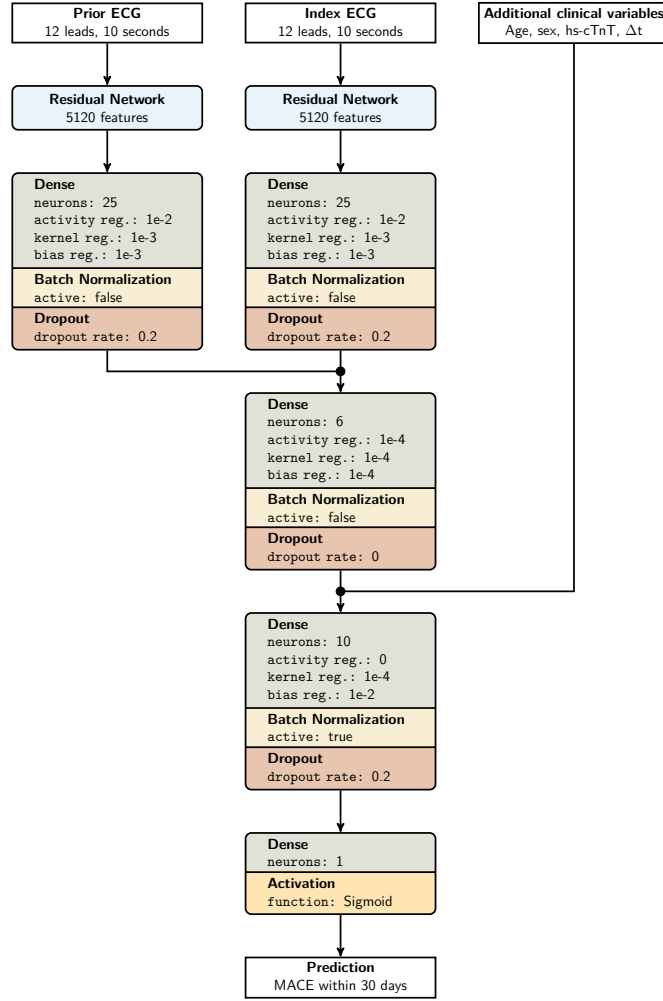


Figure 16: Parameters for and structure of the residual network model using both index and prior ECGs together with additional clinical variables as input.  $\Delta t$  is the time between the two ECGs. The circle denotes concatenation. The learning rate was 0.003 for the first 33 epochs, and 0.00001 for the final 17 epochs.